



Misinformation spreads like a virus. So we built the immune system.

A proposal to stop treating disinformation as a content problem — and start treating it as a public health emergency.

UNESCO YOUTH HACKATHON 2026

Play Your Part: Youth Designing the Future of Media and Information Literacy

LIVE NOW voice-beta-five.vercel.app

STATUS **deployed, not a mockup**

LANGUAGES **24 at launch**

DIAGNOSIS **under 1 second**

COST TO RUN **\$0**

LIVE LINK:

voice-beta-five.vercel.app

SOURCE CODE:

github.com/beastzex/UNESCO_Youth_Hackathon_Submission_2026

DEMO VIDEO:

youtu.be/Tzlr1Jlgyo

For the best experience, view this video on mobile.

OPEN CITIZEN SURVEILLANCE PROTOCOL · © 2026 VOICE INITIATIVE

■ The One-Sentence Version

We turned media literacy from a school subject into a live public-health metric — and shipped it in 24 languages, for free.

■ The Uncomfortable Truth

Every institution fighting disinformation today is organised like a newspaper. A small expert desk. One claim at a time. Corrections published days later.

That is fighting cholera by writing articles about cholera.

The work is not bad. The *shape* is wrong. You cannot beat a transmission problem with a publishing schedule.

Public health solved this a century ago — not with better pamphlets, but with **surveillance networks, case reporting, strain typing, vaccine production and distributed immunisation**. A system where ordinary people are the sensors *and* the delivery mechanism. Where protection is not a personal virtue, but **a number on a map that rises when the community acts**.

VOICE is that system, rebuilt for the information ecosystem.

Five roles. One protocol. A number that tells you whether it is working.

Any citizen is a case reporter	No credentials. No training. No gatekeeper.
Every report is diagnosed in under a second	An open-weight 120-billion-parameter model — not a queue
Rumours are typed into strains	Across language, wording and format — the way pathogens are typed
Counter-content is a two-sentence vaccine	Engineered to survive being forwarded
Protection is measured live	The Herd Immunity Index — per region, right now

And here is the part that matters: this is not a concept, a deck or a mockup. It is **deployed, public and running right now**.

Live link: voice-beta-five.vercel.app — in 24 languages, with zero configuration and zero cost. You can try to break it in the next three minutes.


1 Truth Is Losing a Race It Never Agreed to Run

The six-times problem

Vosoughi, Roy and Aral analysed roughly 126,000 cascades on Twitter (*Science*, 2018). The finding is brutal:

False news reached 1,500 people six times faster than true news. It spread farther, deeper and faster across every category. And it was driven by **humans, not bots.**

Novelty beats correction. That is not a moderation bug to be patched — it is a property of the medium itself. **Any response running on a slower clock than the pathogen loses by definition.**



Five reasons the current model cannot close the gap

Centralised	A desk of 5–20 people is the bottleneck for a population of millions. Throughput can never scale to the threat.
Reactive	Work starts <i>after</i> a claim goes viral. The first two hours — the whole fight — are given away for free.
Long-form	A 1,500-word debunk does not travel through the channels a 15-second clip travels through. Format mismatch is fatal.
English-first	Verification tooling clusters in high-resource languages, leaving communities most dependent on forwarded voice notes structurally undefended.
Unmeasured	No fact-checking organisation on earth can tell you how protected a specific district is <i>today</i> . Without a metric there is no operations — only output.

What the science actually says

The dramatic “backfire effect” has weakened under replication — Wood & Porter (2019) found corrections generally *do* move beliefs. But something more stubborn survived: the **continued influence effect**. Retracted misinformation keeps shaping how people reason *even after they accept the correction*.

Meanwhile, **prebunking works**. Inoculation theory — showing someone a weakened dose of a manipulation *technique* before they meet the real thing — has produced durable, scalable effects in field experiments delivered as short video (Roozenbeek, van der Linden et al., *Science Advances*, 2022).

The evidence points one direction: get there first, and teach the trick — not just the fact.

No deployed platform is organised around that principle. **This one is.**

2 Stop Writing Articles. Start Running an Immune System.

Media and Information Literacy is not a curriculum. It is an immune system.

PUBLIC HEALTH	VOICE
Pathogen	Strain — a clustered false claim
Case report	Submission — one citizen sighting
Diagnosis	AI classification into one of six manipulation techniques
Strain typing	Semantic clustering across rewording, translation and format
Vaccine	A two-sentence plain-language prebunk
Immunisation campaign	Distribution into a named region’s community channels
Herd immunity	A live per-region protection score
Outbreak map	The surveillance radar

Every row above is implemented in code today. Not planned. Not mocked. Running.

Five roles. Anyone can hold one.

01	02	03	04	05
+-----+ SPOTTER -->	+-----+ ANALYST -->	+-----+ VACCINE MAKER -->	+-----+ FIELD WORKER -->	+-----+ RADAR
+-----+ anyone	+-----+ fact-checker	+-----+ educator communicator	+-----+ community leader	+-----+ public official
ingest 60 sec	-> classify under 1 sec	-> synthesize 2 sentences	-> broadcast per region	-> measure live index

Here is the design insight everyone else misses. Each role is a *small, separate, finishable job*. A student is a Spotter in sixty seconds. A teacher is a Vaccine Maker in five minutes. A community leader is a Field Worker with the phone already in their hand.

Nobody has to be an expert at the whole chain — which is exactly why the chain can scale to a population when a monolithic “fact-checker” role never could.

The number that changes everything

$$\text{Herd Immunity}(r) = \frac{\text{strains with a vaccine delivered to } r}{\text{strains active in } r} \times 100$$

70–100%

Stable community inoculation

35–69%

Inoculation in progress

0–34%

Critical vector outbreak

This is the most important thing in this document.

Every other MIL intervention produces *outputs* — lessons delivered, articles published, workshops run. Activity metrics. Attendance sheets. This produces an **outcome that moves in real time and can be targeted**.

- A ministry opens a map and sees which district sits at 12%.
- A school network watches its own number climb week over week.
- A funder measures a programme against something other than a headcount.

It converts media literacy from advocacy into operations.

You cannot manage what you cannot measure. **We built the measurement.**

3 This Is Not a Pitch Deck. It Is Deployed.

Most proposals describe what *will* be built. **This one describes what is already running.**

CAPABILITY	STATUS
Five role consoles, end to end	• Complete
AI classification into 6 techniques with confidence scoring	• Complete
Semantic strain clustering across rewording and language	• Complete
Herd Immunity Index, live-computed	• Complete
Outbreak radar with per-region dossiers	• Complete
24-language interface — 3,624 translated strings	• Complete
D0MI — embedded AI assistant, MIL-briefed, multilingual	• Complete
Deterministic offline fallback for every AI call	• Complete
Runs with zero configuration, zero cost	• Complete
Full technical documentation — 11 documents	• Complete

Go test it: voice-beta-five.vercel.app. A seeded outbreak — five strains across five regions — is already in progress. Submit a rumour. Triage it. Write a vaccine. Broadcast it. **Watch the index move.** Three minutes, start to finish.

Three engineering calls that make the claims real

Sub-second diagnosis

Classification runs on `openai/gpt-oss-120b` via Groq’s LPU inference. That is the difference between a citizen seeing a verdict *while they are still suspicious* and being told to check back tomorrow. **Latency is not a spec here. It is the whole thesis.**

It works when nothing works

Every AI call has a deterministic fallback — a keyword-ladder classifier and a lexical-overlap strain matcher. The platform degrades; it never dies. **Resilience infrastructure that needs reliable infrastructure is not resilience infrastructure.**

Local-first, cloud-optional

State lives in the browser and mirrors to Postgres when configured. The consequence is radical: **the barrier to running VOICE is zero dollars and zero accounts**. A teacher in any country can open it in a classroom this afternoon.

4 Why Nothing Else Does This

	FACT-CHECK DESKS	PLATFORM MODERATION	MIL CURRICULA	VOICE
Response time	24–72 h	Hours to days	A semester	Under 1 second
Who participates	~10 experts	Platform staff	Enrolled students	Anyone
Output	1,500-word article	A label	A lesson plan	2-sentence dose
Timing	After virality	After a report	Before, generically	Before, per strain
Languages	Usually 1–3	Uneven	Local only	24 at launch
Cross-platform	—	One platform	—	Channel-agnostic
Measures protection	✗	✗	Survey proxies	• Live index
Cost to operate	Salaried desk	Corporate scale	Institutional	\$0 to start

Look at the last two rows. Nothing else in this space produces a live protection metric. Nothing else can be stood up by a school or a youth club, for free, this week.

5 The Math of Going Wide

***These are projections from stated assumptions — not measured results.** VOICE has not yet run a field pilot; establishing these numbers empirically is Phase 2. We would rather show our arithmetic than quote a number we have not earned.*

Assumptions. One trained Field Worker maintains ~4 community groups of ~150 members. A vaccine reaches ~40% of a group within 24 hours of posting. Prebunk retention follows the published inoculation literature — meaningful effect persisting on the order of weeks.

DEPLOYMENT	FIELD WORKERS	GROUPS	PEOPLE REACHED, PER STRAIN
One school club	20	80	~4,800
One district network	250	1,000	~60,000
One national youth programme	5,000	20,000	~1,200,000

This scales the way vaccination campaigns scale — **through people already trusted inside the channel**, not through paid reach. That is precisely why we chose the epidemiological model over a media model.

The asymmetry we are exploiting

A strain-specific prebunk costs one educator **five minutes**. The misinformation cost the adversary **five minutes**.

For the first time, defence and offence run on the same clock.

6 The Road From Here

Phase 1 • Harden it — months 0–3

WORK	OUTCOME
Route all inference through server handlers	Live model classification across the entire pipeline
Authentication and per-role authorisation	The five roles become real permission boundaries
Row-level security, audit trail, moderation queue	Safe to deploy with real citizen data
Localise the five role consoles	Working surfaces reach the 24 languages the landing page already has
RTL layout support	Arabic rendered correctly — not merely present
Encoding repair across the dictionaries	Every language displays cleanly

Phase 2 • Prove it — months 3–9

- **Three pilot sites**, deliberately unlike each other: one high-resource urban school network, one multilingual rural district, one youth organisation in a language current fact-checking infrastructure ignores.
- **Test the index against reality**. The essential scientific question is whether the Herd Immunity Index correlates with *measured belief resistance*. Paired pre/post instruments against a matched control region.
- **Publish the methodology openly — including negative results**.

This phase matters more than any other. If the index does not predict real resilience, it gets redesigned. We would rather learn that in a pilot than after scale.

Phase 3 • Federate — months 9–24

TRACK	WHAT IT UNLOCKS
Strain exchange protocol	Independent deployments share strain intelligence across borders. A deepfake technique detected in one country becomes a prebunk template in six others within hours .
Messaging-channel integration	Vaccines delivered directly into WhatsApp, Telegram and Signal — where the pathogen already lives
Embeddings-based clustering	Vector retrieval replaces linear scanning; directories in the tens of thousands
Public API	Newsrooms, election commissions and health ministries consume the outbreak feed
Curriculum bridge	Aligned to the UNESCO MIL Curriculum, so classroom work produces real defensive output , not exercises

Where this ends up

A federated, youth-operated surveillance network in which a manipulation technique first seen anywhere is neutralised everywhere — inside the same news cycle.

7 Built on UNESCO’s Own Blueprint

The Five Laws of MIL, turned into software

LAW	HOW VOICE IMPLEMENTS IT
Information, communication, media and digital content are equal in stature	The taxonomy treats text, image, video, audio and screenshot as peer vectors
Every citizen is a creator of information and has a message	Every citizen is a case reporter — and a potential Vaccine Maker
Information may be used for good or ill; its value is contextual	The classifier reports technique and intent — it teaches mechanism, not verdict
Citizens want to know and understand new information	D0MI answers, in 24 languages, at the moment of curiosity
MIL is a lived, dynamic experience and process	The protocol is a continuous loop — never a completed course

Wider alignment

- **WHO infodemic management** — surveillance, case reporting and response is the WHO’s own infodemic vocabulary. VOICE is a working implementation of it.
- **IPDC** — open source, freely deployable, designed for low-resource communication environments.
- **Global MIL Week** — a live platform participants *operate*, not a demo they watch.
- **SDG 16.10** — public access to verified information as a governance outcome.

Why youth, specifically

Young people are simultaneously **the most exposed cohort to synthetic media and the most fluent in the channels where it spreads**. Every existing response treats them as the population to be *protected*.

VOICE treats them as the network that does the protecting.

That inversion *is* the meaning of *Play Your Part* — and it is encoded directly into the software. **The first role in the protocol requires no credential at all.**

8 How This Could Fail (And What We Did About It)

We would rather say these out loud than have them found.

RISK	SEVERITY	WHAT WE DID
Weaponisation — coordinated actors flood the queue to bury real strains, or report true information as false	High	Analyst confirmation is a mandatory human gate . No submission becomes a strain automatically. Phase 1 adds authentication, rate limiting and an audit trail; Phase 2 adds reputation weighting.
The index gets gamed — distribution marked without real delivery	High	The metric is deliberately self-reported at prototype stage. Phase 2 introduces delivery attestation and spot-audit sampling. A metric that cannot be verified will be treated as an indicator — never a certification.
Model error at scale	Medium	The AI <i>suggests</i> ; a human <i>confirms</i> . Every classification shows its confidence score. Analysts override technique, title and summary before anything commits.
The index does not predict real resilience	High	This is the explicit falsification test of Phase 2. Published either way.
Institutional capture — a state actor runs a deployment to launder its own narrative	Medium	Open source, federated by design, no central authority. The strain exchange protocol carries provenance so consumers can weight sources.
Volunteer burnout	Medium	Roles are deliberately small and separable. Nobody carries the whole chain.

One thing we will never claim: that this platform is a neutral arbiter of truth. It is a coordination and measurement layer for **human judgement**. The Analyst gate is not a bottleneck we intend to remove — **it is the ethical architecture.**

9 What We Are Asking For

Right now, costing nothing: open the live link below and run the loop.

Live link: voice-beta-five.vercel.app

Three minutes. **The argument in this document is either obvious in the product or it is not.**

For the hackathon: judge us on one claim — that **MIL should be measurable infrastructure** — and on whether the Herd Immunity Index is a credible first attempt at that measurement.

NEED	WHY
Three pilot partners	School networks or youth organisations to run Phase 2 — ideally including one non-Latin-script language community
Methodological review	Inoculation-theory researchers to critique and co-design the index validation instrument
Regional MIL expertise	Native speakers to audit the 24 dictionaries and shape culturally-fitted vaccine templates
Modest infrastructure	It runs on free tiers today; pilots need managed hosting and inference quota

We are not asking for a mandate or an endorsement. We are asking for the chance to be tested.

10 The Last Word

Every generation gets an information technology it is not yet immune to. Print gave us a century of religious war before the institutions caught up. Broadcast gave us propaganda states.

We are early in the synthetic-media era — and the institutional response is currently **a newspaper column arriving three days late.**

The thing that beat cholera was never better writing about cholera.

It was **a network of ordinary people who knew what to look for, a system that told them what they had found, and a number that told everyone whether it was working.**

Misinformation spreads like a virus. We built the immune system.

It is running right now, in 24 languages, and it costs nothing to try.

LIVE LINK:

voice-beta-five.vercel.app

SOURCE CODE github.com/beastzex/UNESCO_Youth_Hackathon_Submission_2026

© 2026 VOICE INITIATIVE · OPEN CITIZEN SURVEILLANCE PROTOCOL
MADE WITH LOVE FOR THE UNESCO YOUTH HACKATHON 2026

References

- Vosoughi, S., Roy, D., & Aral, S. (2018). *The spread of true and false news online.* **Science**, 359(6380), 1146–1151.

2. Roozenbeek, J., van der Linden, S., Goldberg, B., Rathje, S., & Lewandowsky, S. (2022). *Psychological inoculation improves resilience against misinformation on social media*. **Science Advances**, 8(34).
 3. Lewandowsky, S., Cook, J., Ecker, U. K. H., et al. (2020). *The Debunking Handbook 2020*.
 4. Wood, T., & Porter, E. (2019). *The elusive backfire effect: Mass attitudes' steadfast factual adherence*. **Political Behavior**, 41, 135–163.
 5. UNESCO. *Media and Information Literacy Curriculum for Educators and Learners; Five Laws of Media and Information Literacy*.
 6. World Health Organization. *Infodemic management: a public health approach*.
-

Repository & Live Platform

Source code github.com/beastzex/UNESCO_Youth_Hackathon_Submission_2026

Live platform voice-beta-five.vercel.app

The repository carries the full source, the technical documentation set, and this proposal.